

wbs

WARWICK BUSINESS SCHOOL
THE UNIVERSITY OF WARWICK

**For the
Change
Makers**

Advanced Programming for Data Science

**Week 10: Neural Network and Review
Information Systems and Management
Warwick Business School**

Working with Jupyter-notebook for Assessment

- Example not template.
- Markdown "text" as your analysis narrative
- Organize your analysis with Markdown "text".

Neural Network and Deep Learning

TensorFlow vs. PyTorch



TensorFlow




Keras

PyTorch
SKORCH

Tensor

- Tensors are similar to NumPy's ndarrays, with the addition being that Tensors can also be used on a GPU to accelerate computing.
- A vector is a one-dimensional or first order tensor and a matrix is a two-dimensional or second order tensor.

```
x = torch.empty(5, 3)  
print(x)
```



First "deep" neural network


```
from tensorflow.keras import Sequential
from tensorflow.keras.layers import Dense
network = Sequential()
network.add(Dense(8, activation='relu', input_dim=4))
network.add(Dense(3, activation='softmax'))
network.compile(loss='categorical_crossentropy',
optimizer='adam', metrics=['accuracy'])
network.fit(train, labels, epochs=100, batch_size=5)
```

Creating the network

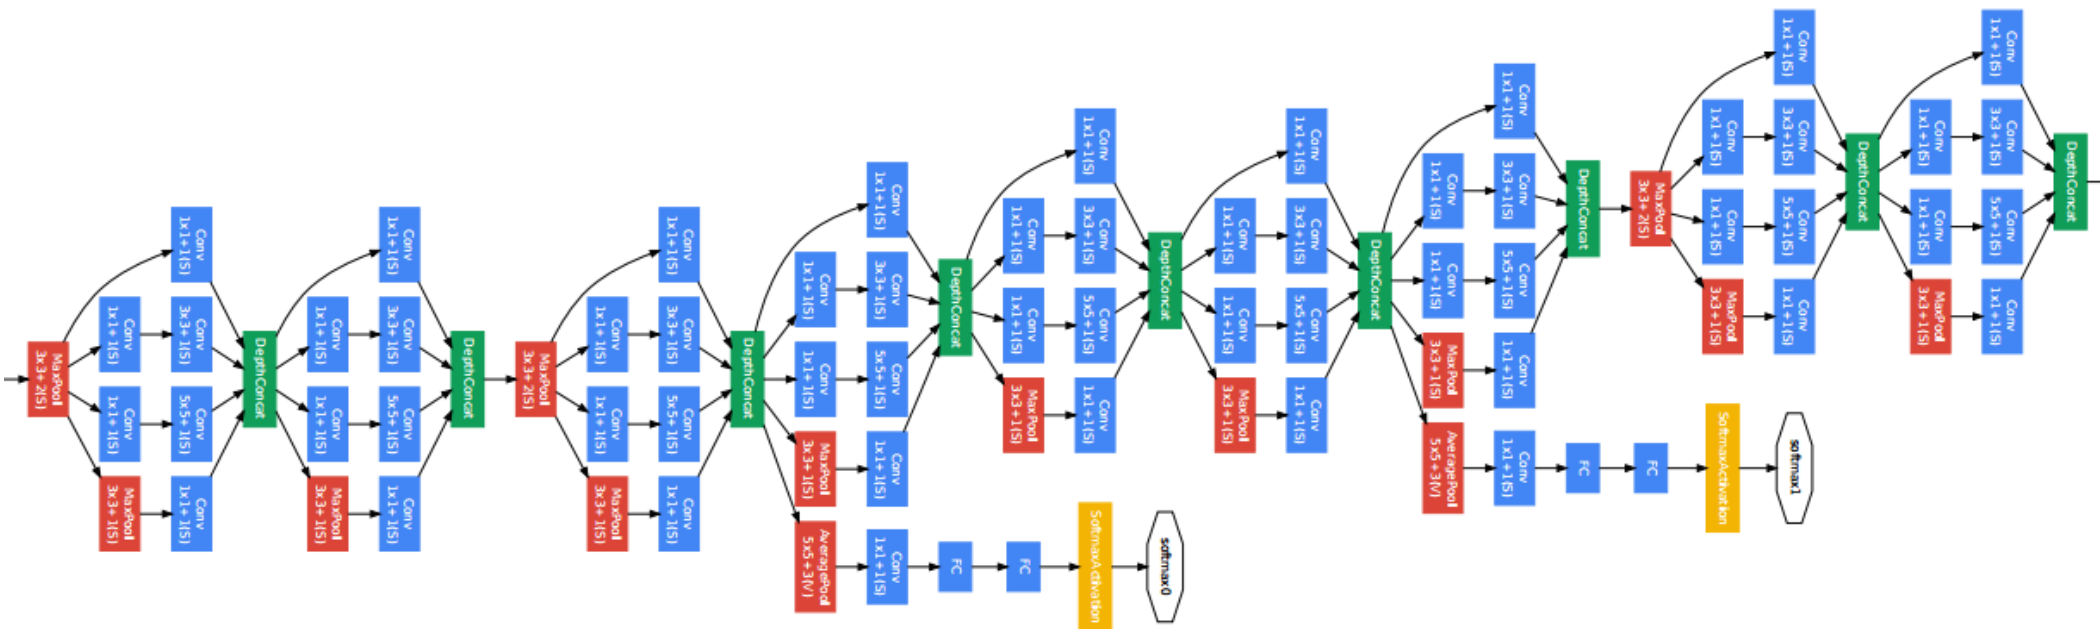
```
network = Sequential()  
network.add(Dense(8, activation='relu', input_dim=4))  
network.add(Dense(3, activation='softmax'))  
network.compile(loss='categorical_crossentropy', optimizer='adam',  
metrics=['accuracy'])  
network.fit(train, labels, epochs=100, batch_size=5)
```


network = Sequential()

- Initialize a network
- A deep neural network is basically a series of layers of "neurons".
- A 3-layer fully connected network.
- Known as multilayer perceptron (MLP) or vanilla neural network.




<https://towardsdatascience.com/coding-neural-network-forward-propagation-and-backpropagation-ccf8cf369f76>





A close look of neural network



Neural Network

$$y = fnn(x) = f_3(f_2(f_1(x)))$$

$$f_l(z) = \sigma_l(Wl z + bl)$$



`network.add(Dense(8, activation='relu', input_dim=4))`

- Creating a deep neural network is like building lego set.
- You can add different layers sequentially.
- We add our first layer, which is a dense layer, with a 4 dimensional input, 8 dimensional output, and an activation function of "relu".
 - Dense means it is a densely-connected (also called "fully-connected") neural layer.
 - It takes input from our original training set, which is a 4-dimensional data.
 - It outputs an 8-dimensional result for next layer.

Activation function

- Activation function is the function to determine the final outcome of the neuron.
- It decides if the neuron is activated or not.
- Therefore, activation function is a non-linear function that generally transforms the linear operation output into a value between 0 to 1.
- Why?
 - Stacking of linear functions is still a linear function, which defeat the purpose of multiple layers.

Common activation functions




`network.add(Dense(3, activation='softmax'))`

- We further add the second layer, which is also the last layer, the output layer.
 - It is still a dense layer.
 - From the second layer, you don't need to specify the input shape as the model will automatically determine it from the previous layer.
 - Since this is the last layer, we aim to produce 3 groups of results. Thus the output dimension is 3.
 - We use an activation function 'softmax', which is normally only used in the last layer for classification problems.

`network.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])`

- Overall flow



Loss function

- Loss function is a function used to assess how well does the model perform in terms of predicting the expected results. Most commonly used include:
 1. MSE (Mean Square Error)
 2. Binary_crossentropy
 3. Categorical_crossentropy

Problem type	Last-layer activation	Loss function
Binary classification	sigmoid	binary_crossentropy
Multi-class, single-label classification	softmax	categorical_crossentropy
Multi-class, multi-label classification	sigmoid	binary_crossentropy
Regression to arbitrary values	None	mse
Regression to values between 0 and 1	sigmoid	mse or binary_crossentropy

Learning: Weight update

For each neuron, $\text{output} = \text{relu}(\text{dot}(W, \text{input}) + b)$. Initially, W and b will be assigned randomly. They will then get updated based on the loss and optimizer.

Gradient-based optimization

Optimizer

- Adam
- Gradient Descent (SGD, Mini Batch, etc.)
- Momentum

Tweaking the network

- Develop a model that does better than a baseline.
 1. Choice of the last-layer activation.
 2. Choice of loss function.
 3. Choice of optimization configuration.
- Scale up: develop a model that overfits
 1. Add layers.
 2. Make your layers bigger.
 3. Train for more epochs.
- Regularize your model and tune your hyperparameters.
- Dropout
- Different architectures.
- Different hyperparameters.
- Iterate feature engineering.

Review

Data Collection

To-do	Tools
Data retrieve via SQL queries	BigQuery, MySQL, etc.
Web Scraping	Requests, BeautifulSoup, etc.
Data access via JSON	json, csv

Data Collection

To-do	Tools
Data retrieve via SQL queries	BigQuery, MySQL, etc.
Web Scraping	Requests, BeautifulSoup, etc.
Data access via JSON	json, csv

Data Visualization

To-do	Tools
Visual inspection	Matplotlib, Seaborn, etc.
Presenting results	Tableau, Google Data Studio

Data Processing

To-do	Tools
Load the dataset	Pandas, csv, etc
Deal with missing data	Pandas: dropna(), fillna() Sklearn imputer.
Wrangling	Pandas: drop(), get_dummies(), cut(), .str.extract() Sklearn.preprocessing: SimpleImputer, StandardScaler, MinMaxScaler, OneHotEncoder, QCA, etc.

Machine Learning

To-do	Tools
Unsupervised Learning: clustering <ul style="list-style-type: none">✓ n_clusters,✓ max_iter,✓ Etc.	sklearn.cluster: KMeans
Supervised Learning: classification <ul style="list-style-type: none">✓ Decision Tree✓ Ensemble✓ criterion, max_depth, etc.	sklearn.tree: DecisionTreeClassifier sklearn.ensemble: RandomForestClassifier, AdaBoostClassifier, StackingClassifier, etc.
Supervised Learning: regression <ul style="list-style-type: none">✓ Linear regression✓ Tree regressor	sklearn.linear_model: LinearRegression sklearn.tree: DecisionTreeRegressor

Machine Learning

To-do	Tools
Model selection ✓ Metrics <ul style="list-style-type: none">• accuracy, recall, precision, f1, etc.• mse, mae, etc. ✓ Cross validation ✓ Grid Search	sklearn.metrics: accuracy_score, recall_score, precision_score, f1_score, etc. sklearn.metrics: mean_absolute_error, mean_squared_error, etc. Sklearn.model_selection: Kfold, StratifiedKfold Sklearn.model_selection: GridSearchCV
Sklearn pipeline	ColumnTransformer, FunctionTransformer, Pipeline, etc.
Deep learning	TensorFlow, PyTorch, Keras, etc.